

Data Integration & Exchange

Guillaume Raschia — Polytech Nantes ; Nantes Université

Last update: December 4, 2023

Introduction

Logic and Database Queries

Answering Query Using Views

Schema Mapping Languages

[Source : L. Libkin - Univ. of Edinburgh, 2015]

[Source : A. Doan, A. Halevy and Z. Ives - "Principles of DI" Series of slides, 2012]

- **Text Book:** A. Doan, A. Halevy and Z. Ives - *Principles of Data Integration*, 2012
- **Text Book:** M. Arenas, P. Barceló, L. Libkin and F. Murlak - *Foundations of Data Exchange*, 2014
- **Chapter In Text Book:** S. Abiteboul et al. - *Web Data Management*, 2011

Introduction

Traditional Approach to Databases

As It Used To Be¹

- A single large repository of data
 - may be distributed across several servers and sites
 - but remains under one single authority
- DBMS takes care of lots of things for you
 - query processing and optimisation
 - concurrency control
 - enforcing database integrity

¹Title of a short film from Clément Gonzalez (2013) - [link to the video](#)

What do we expect from such a system?

1. Data is relatively **clean**: incompleteness is marginal
2. Data is **consistent**: enforced by the DBMS
3. Data is **available**: either on disk or on the local network
4. Queries have a **well-defined semantics**: you know what you pay for
5. Access to data is **controlled**

Traditional Approach to Databases cont'd

This model works well within a single organisation that either

- does not interact much with the outside world, or
- the interaction is heavily controlled by the DB administrator

Homogeneity still dominates but the rules are slightly changing...

What Happens These Days

- Many huge repositories are publicly available
 - In fact many are well-organised databases, e.g., **imdb.com**, the CIA World Factbook, many genome databases, open gov's data platforms, the DBLP server of CS publications, etc etc etc
- Many queries *cannot* be answered using a single source
- Often data from various sources needs to be combined, e.g.
 - company mergers
 - restructuring databases within a single organisation
 - combining data from several private and public sources

Query Answering from Multiple Sources

- Data resides in several **autonomous** databases
- They may have **different structures**, **different access policies** etc
- Our view of the world may be very different from the view of the databases we need to use
- Only **portions of the data** from some database could be available
- That is, the sources **do not conform** to the schema of the database into which the data will be loaded

Heterogeneity is the rule!

ETL tools: Extract-Transform-Load

1. Extract data from multiple sources
2. Transform it so it is clean and compatible with the target schema
3. Load it into a database

The Players

Big ones

IBM, Microsoft, Oracle, SAP – all have their ETL products;

- Microsoft and Oracle offer them with their database products

(Not so) Outsiders

- A few independent vendors, e.g. Informatica PowerCenter
- (Sort of) local player: Talend (aka. Qlik from 2023), R&D Center in Nantes
- Several open source products exist, e.g. CloverETL

The Gartner Magic Quadrant for Data Integration 2020

Figure 1. Magic Quadrant for Data Integration Tools



Source: Gartner (August 2020)

Focus

- Data profiling
- Data cleaning
- Simple transformations
- Bulk loading
- Latency requirements

What they don't do yet

- nontrivial transformations
- query answering
- But techniques now exist for data integration and for query answering
- They soon will be reflected in products

Data Profiling and Data Cleaning

Data profiling gives the user a view of data

- Samples over large tables
- Statistics (how many different values etc)
- Graphical tools for exploring the database

Cleaning

- Same properties may have different names
 - e.g. `Last_Name`, `LName`, `LastName`
- Same data may have different representations
 - e.g. `(0131)555-1111` vs `01315551111`,
 - `George Str.` vs `George Street`
- Some data may be just wrong

- Most transformation rules tend to be **simple**:
 - Copy attribute LName to Last_Name
 - Set age to be `current year - DOB`
- Heavy emphasis on industry specific formats (MS Word, Excel, PDF, UN/EDIFACT, etc)
- Little to do with the **general tasks of data integration**

Integration and Exchange

Integration

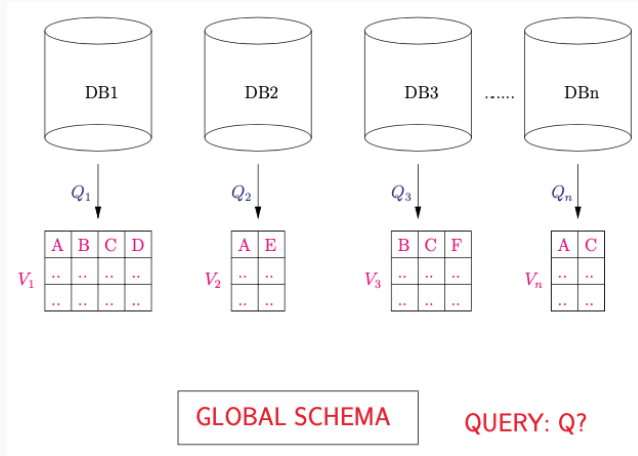
answer queries using multiple sources

- virtual approach, or
- materialization

Exchange

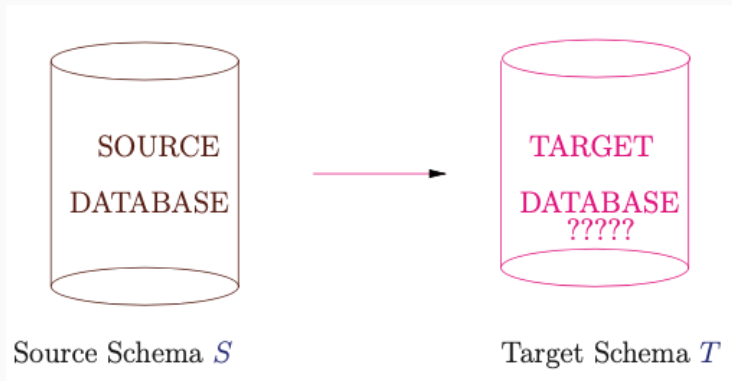
transfer data between two legacy database schemas

Data Integration: Scenario



Answer to Q is obtained by querying the views V_1, V_2, \dots, V_n

Flip Coin: Data Exchange



Query over the target schema: Q

How to answer Q so that the answer is consistent with the **source** database?

What Changes?

- no clear notion of an answer to a query
- data is not clean: incomplete, inconsistent
- data may not even exist (virtual integration)

Our goal

Study the main concepts and techniques for creating and querying integrated/exchanged data

Logic and Database Queries

Questions

How does a data integration system decide which sources are relevant to a query? Which are redundant? How to combine multiple sources to answer a query?

Answer

By reasoning about the content of data sources

- Data sources are described by queries, by views

This section describes the fundamental tools for **manipulating query expressions** and **reasoning** about them.

- Review of basic database concepts
 - Query unfolding
 - Query containment

Relational Terminology

Relational schema

table/relation, attribute/column/field

Relation instance

set (or multi-set) of tuples/rows/records

Integrity constraints

key, foreign key, IND and FD, TGD, EGD

Relational Languages

SQL, RA, RC, Datalog (Rule-based)

Embedded Dependencies (ED)

A fragment of the First-Order logic (FO)

$$\forall \vec{x}, \vec{y} : \phi(\vec{x}, \vec{y}) \rightarrow \exists \vec{z} : \psi(\vec{x}, \vec{z}),$$

ϕ is a possibly empty and ψ is an non-empty conjunction of relational and equality atoms

Two flavors

Tuple Generating Dependency (TGD)

Only relational atoms in ψ

Equality Generating Dependency (EGD)

All atoms in ψ are equalities

 Express FD, Key, IND, and Foreign Key with ED

Conjunctive Queries: CQ

- As logical statements: RC formula with \exists, \wedge only
 - may have interpreted (comparison) predicates

$$\{d, r \mid \exists t, y, s : \text{Movie}(t, d, y) \wedge \text{Rating}(t, r, s) \wedge s > 4\}$$

- As an algebraic expression: Select-Project-Join query (σ , π , and \bowtie only)

$$\pi_{\text{director, reviewer}}(\sigma_{\text{stars} > 4}(\text{Movie} \bowtie \text{Rating}))$$

Conjunctive Queries cont'd

- Rule-based language

$$Q(d, r) \text{ :- Movie}(t, d, y), \text{Rating}(t, r, s), s > 4$$

- $Q(d, r)$ is the **head** of the rule
- $\text{Movie}(t, d, y), \text{Rating}(t, r, s), s > 4$ is the **body**
- an atom like $\text{Movie}(t, d, y)$ is called a **subgoal**
- conjunctions (\wedge) are replaced by commas
- head variables (d, r) are **distinguished**
- variables that occur in the body but not in the head (t, y , and s) are assumed to be existentially quantified (implicit \exists)
- essentially logic programming notation, without functions

Conjunctive Queries with Negation: CQ^-

SPJ + Set Difference (\setminus) queries

$$Q(d, r) :- \text{Movie}(t, d, y), \text{Rating}(t, r, s), s > 4, \neg \text{Rating}(t', r, s'), s' < 2$$

Safe rule

- Every distinguished variable (d, r) must appear in a **positive subgoal**

SPJD + Set Union (\cup) queries \equiv RA \equiv RC

- Multiple rules with the same head

$$Q(d, r) :- \text{Movie}(t, d, y), \text{Rating}(t, r, s), s > 4, \neg \text{Rating}(t', r, s'), s' < 2$$
$$Q(d, r) :- \text{Movie}(t, d, y), \text{Rating}(t, r, s), \text{Reviewer}(r, a), a = \text{'Nantes'}$$

- ✓ Review of basic database concepts
 - Query unfolding
 - Query containment

Query Unfolding

- Query composition is an important mechanism for writing complex queries
 - Build query from views in a bottom up fashion
- Query unfolding “unwinds” query composition
- Important for:
 - Comparing between queries expressed with views
 - Query optimization (to examine all possible join orders)
 - Unfolding may even discover that the composition of two satisfiable queries is unsatisfiable! (✎ find such an example)

Query Unfolding: Example

Database schema: Flight(source, destination) and Hub(city)

$$Q_1(x, y) :- \text{Flight}(x, z), \text{Hub}(z), \text{Flight}(z, y)$$

$$Q_2(x, y) :- \text{Hub}(z), \text{Flight}(z, x), \text{Flight}(x, y)$$

$$Q_3(x, z) :- Q_1(x, y), Q_2(y, z)$$

The unfolding of Q_3 is:

$$Q'_3(x, z) :- \text{Flight}(x, u), \text{Hub}(u), \text{Flight}(u, y), \text{Hub}(w), \text{Flight}(w, y), \text{Flight}(y, z)$$

- mapping $f = \{x/x, y/y, z/u\}$ ² in Q_1 and $g = \{x/y, y/z, z/w\}$ in Q_2

² z/u denotes $f(z) = u$ and it roughly means “replace z by u ”.

Query Unfolding Algorithm

1. Find a subgoal $P(x_1, \dots, x_n)$ such that P is defined by a rule r
2. Unify $P(x_1, \dots, x_n)$ with the head of r
3. Replace $P(x_1, \dots, x_n)$ with the result of applying the unifier to the subgoals of r (use fresh variables for the existential variables of r)
4. Iterate until no unification can be found
5. If P is defined by a union of rules r_1, \dots, r_m , create m rules of the main query, one for each of the r 's

Query Unfolding: Summary

- Unfolding does not necessarily create a more efficient query!
 - Just let's the optimizer explore more evaluation strategies
 - Unfolding is the opposite of rewriting queries using views (see later on)
- The size of the resulting query can grow exponentially (✎ show how)

- ✓ Review of basic database concepts
- ✓ Query unfolding
 - Query containment

Query Containment: Motivation

Intuitively, the unfolding of Q_3 is **equivalent** to Q'_3 :

$$Q'_3(x, z) :- \text{Flight}(x, u), \text{Hub}(u), \text{Flight}(u, y), \text{Hub}(w), \text{Flight}(w, y), \text{Flight}(y, z)$$

$$Q''_3(x, z) :- \text{Flight}(x, u), \text{Hub}(u), \text{Flight}(u, y), \text{Flight}(y, z)$$

since, mainly, $\text{Hub}(w)$ and $\text{Flight}(w, y)$ can be unified with the same tuples than $\text{Hub}(u)$ and $\text{Flight}(u, y)$

- How can we justify this intuition formally?

Query Containment: Motivation cont'd

Furthermore, the query Q_4 that requires going through two hubs is **contained** in Q_3'' (and Q_3' as well)

$$Q_3''(x, z) :- \text{Flight}(x, u), \text{Hub}(u), \text{Flight}(u, y), \text{Flight}(y, z)$$

$$Q_4(x, z) :- \text{Flight}(x, u), \text{Hub}(u), \text{Flight}(u, y), \text{Hub}(y), \text{Flight}(y, z)$$

- We need algorithms to detect these relationships

Query Containment and Equivalence: Definitions

Definition (Query Containment)

Query Q_1 is contained in query Q_2 if for every database D , $Q_1(D) \subseteq Q_2(D)$

Definition (Query Equivalence)

Query Q_1 is equivalent to query Q_2 if for every database D :

$$Q_1(D) \subseteq Q_2(D) \quad \text{and} \quad Q_2(D) \subseteq Q_1(D)$$

Note

Containment and equivalence are properties of the queries, not the database!

$$Q_1(x, z) :- P(x, y, z)$$

$$Q_2(x, z) :- P(x, \textcolor{brown}{x}, z)$$

$$Q_2 \sqsubseteq Q_1$$

$$Q_1(x, y) :- P(x, z), P(z, y)$$

$$Q_2(x, y) :- P(x, z), P(z, y), P(x, w)$$

$$Q_1 \sqsubseteq Q_2 \quad \text{and} \quad Q_2 \sqsubseteq Q_1$$

Why Do We Need It?

- When sources are described as **views**, we use containment to compare among them
- If we can remove subgoals—joins—from a query, its evaluation is more efficient
- Actually, containment arises everywhere...

Reconsidering the Example

Database schema: Flight(src, dest) and Hub(city)

Views:

$$Q_1(x, y) :- \text{Flight}(x, z), \text{Hub}(z), \text{Flight}(z, y)$$

$$Q_2(x, y) :- \text{Hub}(z), \text{Flight}(z, x), \text{Flight}(x, y)$$

And query: $Q_3(x, z) :- Q_1(x, y), Q_2(y, z)$

The unfolding of Q_3 is:

$$Q'_3(x, z) :- \text{Flight}(x, u), \text{Hub}(u), \text{Flight}(u, y), \text{Hub}(w), \text{Flight}(w, y), \text{Flight}(y, z)$$

Remove Redundant Subgoals

$Q'_3(x, z) :- \text{Flight}(x, u), \text{Hub}(u), \text{Flight}(u, y), \text{Hub}(w), \text{Flight}(w, y), \text{Flight}(y, z)$

$Q''_3(x, z) :- \text{Flight}(x, u), \text{Hub}(u), \text{Flight}(u, y), \text{Flight}(y, z)$

Is Q''_3 truly equivalent to Q'_3 ?

\Leftrightarrow Do both $Q''_3 \sqsubseteq Q'_3$ and $Q'_3 \sqsubseteq Q''_3$ hold?

$$Q(\vec{x}) :- G_1(\vec{x}_1), \dots, G_n(\vec{x}_n)$$

- Assume there is no interpreted predicates (\leq, \neq) nor negative subgoal (\neg)
...at least for now

Recall semantics

If f maps the body subgoals to tuples in D , then $f(\vec{x})$ is an answer

Containment Mapping

$$Q_1(\vec{x}) :- R_1(\vec{x}_1), \dots, R_n(\vec{x}_n)$$

$$Q_2(\vec{y}) :- S_1(\vec{y}_1), \dots, S_m(\vec{y}_m)$$

Definition (Containment Mapping, aka. Homomorphism)

$$f : \text{Variables}(Q_1) \rightarrow \text{Variables}(Q_2)$$

is a containment mapping from Q_1 to Q_2 if:

$$f(R_i(\vec{x}_i)) \in \text{Body}(Q_2), \quad \forall i \in \llbracket 1, n \rrbracket \tag{1}$$

$$\text{and} \quad f(\vec{x}) = \vec{y} \tag{2}$$

Containment Mapping: Example

$Q'_3(x, z) :- \text{Flight}(x, u), \text{Hub}(u), \text{Flight}(u, y), \text{Hub}(w), \text{Flight}(w, y), \text{Flight}(y, z)$

$Q''_3(x, z) :- \text{Flight}(x, u), \text{Hub}(u), \text{Flight}(u, y), \text{Flight}(y, z)$

$f = \{x/x, y/y, u/u, w/u, z/z\}$ is a containment mapping from Q'_3 to Q''_3

1. subgoals: $\text{Hub}(w) \xrightarrow{f} \text{Hub}(u)$, and $\text{Flight}(w, y) \xrightarrow{f} \text{Flight}(u, y)$, both in $\text{Body}(Q''_3)$
2. head variables: $f(x) = x$ and $f(z) = z$

Containment Mapping: Main Result

Theorem [Chandra and Merlin, 1977]

Q_1 contains Q_2 (denoted $Q_2 \sqsubseteq Q_1$) if and only if there is a containment mapping f from Q_1 to Q_2

Deciding whether Q_1 contains Q_2 is NP-complete

Sketch of the Proof

$$Q_1(\vec{x}) :- R_1(\vec{x}_1), \dots, R_n(\vec{x}_n)$$

$$Q_2(\vec{y}) :- S_1(\vec{y}_1), \dots, S_m(\vec{y}_m)$$

The ‘if’ direction

Assume f exists: $f(R_i(\vec{x}_i)) \in \text{Body}(Q_2)$ and $f(\vec{x}) = \vec{y}$

Let t be an answer to Q_2 over database D ; then there is a unification g from the variables of Q_2 to D

Hence, $g \circ f$ is a unification from $\text{Variables}(Q_1)$ to D and t is also an answer to Q_1 #

The 'only-if' direction

Assume the containment $Q_2 \sqsubseteq Q_1$ holds;

Consider the **frozen database** D' of Q_2 :

- Variables of Q_2 are constants in D'

The unification from Q_1 to D' gives the containment mapping!

$$Q(x, z) \text{ :- Flight}(x, u), \text{Hub}(u), \text{Flight}(u, y), \text{Flight}(y, z)$$

Frozen DB for $Q(x, z)$

Flight = $\{(x, u), (u, y), (y, z)\}$

Hub = $\{(u)\}$

Two Views of This Result

1. **Variable mapping:**

- a condition on variable mappings that guarantees containment

2. **Representative (canonical) databases:**

- a (single) database that would offer a counter-example if any

Containment results typically fall into **one of these two classes**

Union of Conjunctive Queries

$$r_{1.1} : Q_1(x, y) :- \text{Flight}(x, z), \text{Flight}(z, y)$$

$$r_{1.2} : Q_1(x, y) :- \text{Flight}(x, z), \text{Flight}(y, z), \text{Hub}(z)$$

and

$$Q_2(x, y) :- \text{Flight}(x, z), \text{Flight}(z, y), \text{Hub}(z)$$

One can observe $Q_2 \sqsubseteq r_{1.1}$

Theorem

A CQ is contained in a union of CQ's iff it is contained in **one** of the conjunctive rules

Corollary

Containment is still NP-complete!

CQ's with Interpreted Predicates

A tweak on containment mappings provides a **sufficient** condition

$$Q_1(\vec{x}) :- R_1(\vec{x}_1), \dots, R_n(\vec{x}_n), C_1$$

$$Q_2(\vec{y}) :- S_1(\vec{y}_1), \dots, S_m(\vec{y}_m), C_2$$

Containment mapping revisited

$f : \text{Variables}(Q_1) \rightarrow \text{Variables}(Q_2)$, with

$$f(R_i(\vec{x}_i)) \in \text{Body}(Q_2), \quad \forall i \in \llbracket 1, n \rrbracket \quad (3)$$

$$f(\vec{x}) = \vec{y} \quad (4)$$

$$\text{and} \quad C_2 \models f(C_1) \quad (5)$$

Containment Mapping of CQ with IP: Example

$Q_1(x, y) :- \text{Flight}(x, z), \text{Flight}(z, y), \text{Population}(z, p), p \leq 500,000$

$Q_2(u, v) :- \text{Flight}(u, w), \text{Flight}(w, v), \text{Hub}(w), \text{Population}(w, s), s \leq 100,000$

Building the mapping f from Q_1 to Q_2 :

- $x \xrightarrow{f} u, y \xrightarrow{f} v, z \xrightarrow{f} w$ and $p \xrightarrow{f} s$, that one can denote $f = \{x/u, y/v, z/w, p/s\}$
- $s \leq 100,000 \models s \leq 500,000$

Containment Mapping: Not a Necessary Condition

$$Q_1(x, y) :- R(x, y), S(u, v), u \leq v$$

$$Q_2(x, y) :- R(x, y), S(u, v), S(v, u)$$

No containment mapping from Q_1 to Q_2 :

- x/x and y/y for the head variables
- u/u and v/v to map the S subgoal
- but $\perp \not\models u \leq v$

Anyway, we have $Q_2 \sqsubseteq Q_1$!

$$Q_1(x, y) :- R(x, y), S(u, v), u \leq v$$

$$Q_2(x, y) :- R(x, y), S(u, v), S(v, u)$$

We consider an equivalent rewriting of Q_2 :

$$Q'_2(x, y) :- R(x, y), S(u, v), S(v, u), u \leq v$$

$$Q'_2(x, y) :- R(x, y), S(u, v), S(v, u), u > v$$

Q'_2 rules are the **refinements** of Q_2

$$Q_1(x, y) :- R(x, y), S(u, v), u \leq v$$

$$r_{2.1} : Q'_2(x, y) :- R(x, y), S(u, v), S(v, u), u \leq v$$

$$r_{2.2} : Q'_2(x, y) :- R(x, y), S(u, v), S(v, u), u > v$$

Two containment mappings can then be defined:

1. from Q_1 to $r_{2.1}$: $\{x/x, y/y, u/u, v/v\}$ with $u \leq v \models u \leq v$, and
2. from Q_1 to $r_{2.2}$: $\{x/x, y/y, u/v, v/u\}$ with $u > v \models v \leq u$

Constructing Query Refinements

1. Consider all **complete orderings** of the variables and constants in the query
2. For each complete ordering, create a conjunctive query
3. The result is the union of conjunctive queries

Complete Ordering

- Given
 - a conjunction C of interpreted atoms over
 - a set of variables x_1, \dots, x_n , and
 - a set of constants a_1, \dots, a_m
- C_T is a **complete ordering** if:
 - $C_T \models C$, and
 - $\forall \vartheta_1, \vartheta_2 \in \{x_1, \dots, x_n, a_1, \dots, a_m\}$

$$C_T \models \vartheta_1 < \vartheta_2 \quad \text{or} \quad C_T \models \vartheta_1 > \vartheta_2 \quad \text{or} \quad C_T \models \vartheta_1 = \vartheta_2$$

$$Q_1(\vec{x}) :- R_1(\vec{x}_1), \dots, R_n(\vec{x}_n), C_1$$

Let C_T be a complete ordering of C_1 ; then

$$Q'_1(\vec{x}) :- R_1(\vec{x}_1), \dots, R_n(\vec{x}_n), C_T$$

is a refinement of Q_1

Theorem [Klug, 88; van der Meyden, 92]

Q_1 contains Q_2 if and only if there is a containment mapping from Q_1 to **every refinement** of Q_2

Deciding whether Q_1 contains Q_2 is Σ_2^p -complete

What about Negation?

$$Q_1(\vec{x}) :- R_1(\vec{x}_1), \dots, R_n(\vec{x}_n), \neg S_1(\vec{y}_1), \dots, \neg S_m(\vec{y}_m)$$

$$Q_2(\vec{w}) :- T_1(\vec{w}_1), \dots, T_k(\vec{w}_k), \neg U_1(\vec{z}_1), \dots, \neg U_\ell(\vec{z}_\ell)$$

Queries are assumed to be **safe**:

- every head variable appears in a positive subgoal in the body

Containment mapping revisited

Map negative subgoals in Q_1 to negative subgoals in Q_2

This is a **sufficient condition**, but not a necessary one

Bag Semantics

Flight =

Source	Destination	Departure Time
Nantes	Paris	8am
Nantes	Paris	10am
Paris	Lyon	1pm

$$Q(x, y) :- \text{Flight}(x, z, w), \text{Flight}(z, y, t)$$

- Answers under set semantics: {(Nantes, Lyon)}
- Answers under **bag semantics**: {(Nantes, Lyon), (Nantes, Lyon)}

Equivalence Theorem

Q_1 is equivalent to Q_2 if and only if there is a 1-1 containment mapping

Trivial example of non-equivalence

$$Q_1(x) :- P(x)$$

$$Q_2(x) :- P(x), P(x)$$

What about query containment?

Answering Query Using Views

Motivating Example

Movie(mID, title, year, genre)

Director(mID, dName)

Actor(mID, aName)

$Q(t, y, d) :- \text{Movie}(i, t, y, g), y \geq 1950, g = \text{comedy}, \text{Director}(i, d), \text{Actor}(i, d)$

$V_1(t, y, d) :- \text{Movie}(i, t, y, g), y \geq 1940, g = \text{comedy}, \text{Director}(i, d), \text{Actor}(i, d)$

- Obviously $Q \sqsubseteq V_1$, hence $Q_0(t, y, d) :- V_1(t, y, d), y \geq 1950$
- Containment is enough to show that V_1 can be used to answer Q

Motivating Example cont'd

$Q(t, y, d) :- \text{Movie}(i, t, y, g), y \geq 1950, g = \text{comedy}, \text{Director}(i, d), \text{Actor}(i, d)$

$V_2(i, t, y) :- \text{Movie}(i, t, y, g), y \geq 1950, g = \text{comedy}$

$V_3(i, d) :- \text{Director}(i, d), \text{Actor}(i, d)$

- Containment does not hold,
- but intuitively, V_2 and V_3 are useful for answering Q

How do we express that intuition? Answering queries using views!

- Review of basic database concepts
 - Query unfolding
 - Query containment

Schema Mapping Languages

How to map **Mediated (aka. Global) Schema** to **Source Schemas**?

Semantic Heterogeneity

- Difference in:
 - Naming of schema elements
 - Organization of tables
 - Coverage and detail of schema
 - Data-level representation: John Doe vs. J. Doe
- Why?
 - schema probably designed for different apps, by different people

Schema Heterogeneity by Example

Mediated Schema

Movie: title, director, year, genre

Actors: title, name

Plays: movie, location, startTime

Reviews: title, rating, description



logic

Sources

S1

Movie(mID, title)

Actor(aID, firstName, lastName, nationality, yearOfBirth)

ActorPlays(aID, mID)

MovieDetails(mID, director, genre, year)

S2

Cinemas(place, movie, start)

S3

NYCCinemas(name, title, startTime)

S5

MovieGenres(title, genre)

S6

MovieDirectors(title, dir)

S7

MovieYears(title, year)

S4


Reviews(title, date, grade, review)

Principles of Schema Mapping

Schema mapping describes the relation between:




or between:



Semantics of Schema Mapping

Formally, schema mapping states: which instances of the mediated schema are **consistent with the current instances** of the data sources



$I(G) \times (I(S_i))_{1 \leq i \leq n}$: the set of **possible instances** of the schema $(G, (S_i)_{1 \leq i \leq n})$

$$M \subseteq I(G) \times I(S_1) \times \dots \times I(S_n)$$

Possible Instances of the Mediated Schema: Simple Example

- Source 1: (**Director**, **Title**, **Year**) with tuples
 - {(Allen, Manhattan, 1979),
 - (Coppola, GodFather, 1972)}
- Mediated schema: (**Title**, **Year**)
 - Simple projection of Source 1
 - Only **one possible instance**: {(Manhattan, 1979), (GodFather, 1972)}

Possible Instances of the Mediated Schema: Second Example

- Source 1: (**Title**, **Year**) with tuples
 - {(Manhattan, 1979), (GodFather, 1972)}
- Mediated schema: (**Director**, **Title**, **Year**)
 - Possible instance 1: {(Allen, Manhattan, 1979), (Coppola, GodFather, 1972)}
 - Possible instance 2: {(Alice, Manhattan, 1979), (Bob, GodFather, 1972)}
- Why it is so important to know?
 - This matters **at query time** (see next slide)

Answering Queries over Possible Instances of Mediated Schema


Mediated schema: (**Director**, **Title**, **Year**)

- Possible instance 1: {(Allen, Manhattan, 1979), (Coppola, GodFather, 1972)}
- Possible instance 2: {(Alice, Manhattan, 1979), (Bob, GodFather, 1972)}
- Query Q_1 : return all years of movies
 - Answer: (1979, 1972) are **certain answers**
- Query Q_2 : return all directors
 - No certain answer because no directors appear in all possible instances of the mediated schema

Certain Answers Make it Formal

An answer is **certain** if it is **true in every instance** of the mediated schema that is consistent with:

1. the instances of the sources, and
2. the mapping M

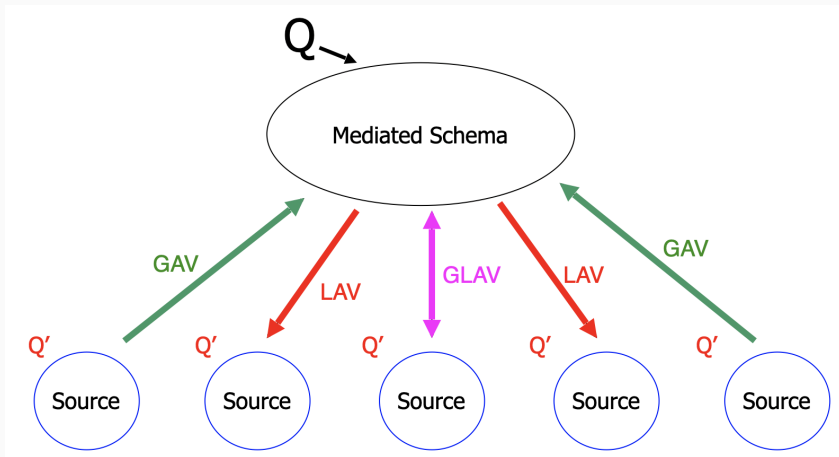


$$\text{certain}(Q, s_1, \dots, s_n) = \bigcap Q(g), \quad \forall g \text{ such that } (g, s_1, \dots, s_n) \in M$$

Desiderata from Source Description Languages

- Flexibility:
 - Should be able to express relationships between real schemata
- Efficient reformulation:
 - Computational complexity of reformulation and finding answers
- Easy update:
 - Should be easy to add and delete sources

Languages for Schema Mapping



Mediated schema is defined as a **set of views over the data sources**

- $G: \text{Movie}(\text{title}, \text{director}, \text{year}, \text{genre})$
- S_1 :
 - $\text{Movie}(\text{mID}, \text{title})$
 - $\text{MovieDetails}(\text{mDI}, \text{director}, \text{genre}, \text{year})$

a **sound** GAV setting:

$\text{Movie}(\text{title}, \text{director}, \text{year}, \text{genre}) \supseteq$
 $S_1.\text{Movie}(\text{mID}, \text{title}), S_1.\text{MovieDetails}(\text{mID}, \text{director}, \text{genre}, \text{year})$

A set of expressions of the form:

$$G_i(\vec{x}) \supseteq Q_i(\vec{S}) \quad \text{or} \quad G_i(\vec{x}) = Q_i(\vec{S})$$

resp. for sound (OWA) or exact (CWA) setting

- G_i is a relation in the mediated schema
- $Q_i(\vec{S})$ is a query over source relations

- $\text{Movie}(\text{title}, \text{director}, \text{year}, \text{genre}) \supseteq$
 $S_1.\text{Movie}(\text{mID}, \text{title}), S_1.\text{MovieDetails}(\text{mID}, \text{director}, \text{genre}, \text{year})$
- $\text{Movie}(\text{title}, \text{director}, \text{year}, \text{genre}) \supseteq$
 $S_5.\text{MovieGenres}(\text{title}, \text{genre}), S_6.\text{MovieDirectors}(\text{title}, \text{director}),$
 $S_7.\text{MovieYear}(\text{title}, \text{year})$

Remind global schema `Plays(movie, location, starTime)`

- $\text{Plays}(m, \ell, s) \supseteq S_2.\text{Cinemas}(\ell, m, s)$
- $\text{Plays}(m, \ell, s) \supseteq S_3.\text{NYCCinemas}(\ell, m, s)$

Reformulation in GAV

- Given a query Q on the mediated schema G ;
 - Return the **best query possible** Q' on the data sources S_1, \dots, S_n
 - Global views **unfolding** is the swiss knife
- Query:

$$Q(t, \ell, s) :- \text{Movie}(t, d, y, \text{'comedy'}), \text{Plays}(t, \ell, s), s \geq 8\text{pm}$$

- Two first GAV rules to unfold the Movie and Plays subgoals of Q :
 - $\text{Movie}(\text{title}, \text{director}, \text{year}, \text{genre}) \supseteq$
 $S_1.\text{Movie}(\text{mID}, \text{title}), S_1.\text{MovieDetails}(\text{mID}, \text{director}, \text{genre}, \text{year})$
 - $\text{Plays}(\text{m}, \ell, s) \supseteq S_2.\text{Cinemas}(\ell, \text{m}, s)$

$$Q(t, \ell, s) :- \text{Movie}(t, d, y, \text{'comedy'}), \text{Plays}(t, \ell, s), s \geq 8\text{pm}$$

becomes

$$Q'(t, \ell, s) :- S_1.\text{Movie}(i, t), S_1.\text{MovieDetails}(i, d, \text{'comedy'}, y), \\ S_2.\text{Cinemas}(\ell, t, s), s \geq 8\text{pm} \quad (6)$$

$$Q(t, \ell, s) :- \text{Movie}(t, d, y, \text{'comedy'}), \text{Plays}(t, \ell, s), s \geq 8\text{pm}$$

becomes

$$Q'(t, \ell, s) :- S_1.\text{Movie}(i, t), S_1.\text{MovieDetails}(i, d, \text{'comedy'}, y), \\ S_3.\text{NYCCinemas}(\ell, t, s), s \geq 8\text{pm} \quad (7)$$

- Recall:
 - $(g, s_1, \dots, s_n) \in M$ if $g \in \mathbf{I}(G)$ is a global database that is consistent with all the source extensions $(s_1, \dots, s_n) \in \mathbf{I}(S_1) \times \dots \times \mathbf{I}(S_n)$
- Then,
 - $G_i(\vec{x}) \supseteq Q_i(\vec{S})$ states the extension of G_i in g is a superset of evaluating Q_i on the sources, or
 - $G_i(\vec{x}) = Q_i(\vec{S})$ states the extension of G_i in g is equal to evaluating Q_i on the sources

Tricky GAV Example

S_8 : stores pairs of (actor, director)

- Then the Actors and Movie global schema can be defined as views like
 - $\text{Actors}(\perp, \text{actor}) \supseteq S_8(\text{actor}, \text{director})$
 - $\text{Movie}(\perp, \text{director}, \perp, \perp) \supseteq S_8(\text{actor}, \text{director})$

Tricky GAV Example cont'd

- GAV setting:
 - $\text{Actors}(\perp, \text{actor}) \supseteq S_8(\text{actor}, \text{director})$
 - $\text{Movie}(\perp, \text{director}, \perp, \perp) \supseteq S_8(\text{actor}, \text{director})$
- Given the S_8 extension: $\{(\text{Keaton}, \text{Allen}), (\text{Pacino}, \text{Coppola})\}$
- We'd get tuples for the mediated schema:
 - $\text{Actors} \supseteq \{(\perp, \text{Keaton}), (\perp, \text{Pacino})\}$
 - $\text{Movie} \supseteq \{(\perp, \text{Allen}, \perp, \perp), (\perp, \text{Coppola}, \perp, \perp)\}$

Tricky GAV Example cont'd

- g extension:
 - $\text{Actors} \supseteq \{(\perp, \text{Keaton}), (\perp, \text{Pacino})\}$
 - $\text{Movie} \supseteq \{(\perp, \text{Allen}, \perp, \perp), (\perp, \text{Coppola}, \perp, \perp)\}$
- Can't answer the query :

$$Q(a, d) :- \text{Actors}(t, a), \text{Movie}(t, d, g, y)$$

Actually, LAV (Local-As-View) setting will solve this problem

- Mediated schema is defined as views over the sources
- Reformulation/unfolding is conceptually easy
 - Polynomial-time reformulation and query answering
- GAV forces everything into the mediated schema's perspective
 - Cannot capture a variety of tabular organizations

Local-As-View (LAV)

Data sources defined as views over the mediated schema

S₅
MovieGenres(title, genre)

S₆
MovieDirectors(title, dir)

S₇
MovieYears(title, year)

Movie: title, director, year, genre

Actors: title, name

Plays: movie, location, startTime

Reviews: title, rating, description

- $S_5.\text{MovieGenres}(t, g) \subseteq \text{Movie}(t, d, y, g)$
- $S_6.\text{MovieDirector}(t, d) \subseteq \text{Movie}(t, d, y, g)$

S8

ActorDirectors(actor, dir)

Movie: title, director, year, genre

Actors: title, name

Plays: movie, location, startTime

Reviews: title, rating, description

- $S_8.\text{ActorDirectors}(a, d) \subseteq \text{Movie}(t, d, y, g), \text{Actor}(t, a), y \geq 1970$

A set of expressions of the form:

$$S_i(\vec{x}) \subseteq Q_i(G) \quad \text{or} \quad S_i(\vec{x}) = Q_i(G)$$

resp. for sound (OWA) or exact (CWA) setting

- S_i is a source relation
- $Q_i(G)$ is a query over mediated schema

- Recall:
 - $(g, s_1, \dots, s_n) \in M$ if $g \in \mathbf{I}(G)$ is a global database that is consistent with all the source extensions $(s_1, \dots, s_n) \in \mathbf{I}(S_1) \times \dots \times \mathbf{I}(S_n)$
- Then,
 - $S_i(\vec{x}) \subseteq Q_i(G)$ states that the result of Q_i over g is a superset of s_i , or
 - $S_i(\vec{x}) = Q_i(G)$ states that the result of Q_i over g is equal to s_i

Unlike GAV, LAV definitions imply a set of possible databases for the mediated schema

- $S_8.\text{ActorDirectors}(a, d) \subseteq \text{Movie}(t, d, y, g), \text{Actor}(t, a)$
- S_8 extension is $\{(Keaton, Allen)\}$
- Two possible databases for the mediated schema are:
 - $\text{Movie}=\{(Manhattan, Allen, 1979, comedy)\}$ and $\text{Actor}=\{(Manhattan, Keaton)\}$
 - $\text{Movie}=\{(Foobar, Allen, 1979, comedy)\}$ and $\text{Actor}=\{(Foobar, Keaton)\}$


In a **sound setting** only, since the source may be incomplete, other tuples may be in the instance of the mediated schema:

- $\text{Movie} = \{(\text{Manhattan}, \text{Allen}, 1979, \text{comedy}), (\text{Leatherheads}, \text{Clooney}, 2008, \text{comedy})\}$
- $\text{Actor} = \{(\text{Manhattan}, \text{Keaton}), (\text{The Godfather}, \text{Keaton})\}$

Certain Answers: A Gentle Reminder

An answer is **certain** if it is **true in every instance** of the mediated schema that is consistent with:

1. the instances of the sources, and
2. the mapping M



$$\text{certain}(Q, s_1, \dots, s_n) = \bigcap Q(g), \quad \forall g \text{ such that } (g, s_1, \dots, s_n) \in M$$

Certain Answers: Example #1

- $S_8.\text{ActorDirectors}(a, d) \subseteq \text{Movie}(t, d, y, g), \text{Actor}(t, a)$
- S_8 extension is $\{(\text{Keaton}, \text{Allen})\}$

$$Q(a, d) :- \text{Movie}(t, d, y, g), \text{Actor}(t, a)$$

- Only one certain answer: (Keaton, Allen)

Certain Answers: Example #2

- $S_9(\text{dir}) \subseteq \text{Director}(t, \text{dir})$, with $s_9 = \{\text{Allen}\}$
- $S_{10}(\text{actor}) \subseteq \text{Actor}(t, \text{actor})$, with $s_{10} = \{\text{Keaton}\}$

$$Q(a, d) :- \text{Director}(i, d), \text{Actor}(i, a)$$

- Under CWA: every possible DB follows the pattern
 - $\text{Director} = \{(x, \text{Allen})\}$, and $\text{Actor} = \{(x, \text{Keaton})\}$
- Then the only exact certain answer is (Keaton, Allen)
- Under OWA: no sound certain answer!

- We're given tuples for the sources (expressed as views)
- We're given a mediated schema (but no tuple)
- We have a query against that mediated schema

This is exactly the problem of **Answering Queries Using Views!**

- Reformulation = **answering queries using views**
- Algorithms work well in practice:
 - Reformulation is not the bottleneck
 - Under some conditions, guaranteed to find **all certain answers**
 - In practice, they typically do
 - LAV expresses **incomplete information**
 - GAV does not: only a single instance of the mediated schema is consistent with sources

S1

Movie(**mID**, title)

Actor(aID, firstName, lastName, nationality, yearOfBirth)

ActorPlays(aID, **mID**)

MovieDetails(**mID**, director, genre, year)

Movie: title, director, year, genre

- If a key is internal to as data source, LAV cannot use it
- So...

Global-and-Local-As-View: The best of all Worlds

- A set of expressions of the form:

$$Q^S(\vec{x}) \subseteq Q^G(\vec{x}) \quad \text{or} \quad Q^S(\vec{x}) = Q^G(\vec{x})$$

resp. for sound (OWA) or exact (CWA) setting

- Q^S is a query over the data sources
- Q^G is a query over the mediated schema

$$S_1.\text{Movie}(i, t), S_1.\text{MovieDetails}(i, d, g, y) \subseteq \text{Movie}(t, d, \text{"comedy"}, y), y \geq 1970$$

- Given a query Q
- Remind the GLAV setting pattern: $Q^S(\vec{x}) \subseteq Q^G(\vec{x})$, then
 1. Find a rewriting Q' using the views Q_1^G, \dots, Q_n^G
 2. Create Q'' by replacing each Q_i^G by the corresponding Q_i^S , then
 3. Unfold Q_1^S, \dots, Q_n^S to get Q'''

Tuple Generating Dependencies (TGD) can be used to specify GLAV expressions

$$(\forall \vec{x}) S_1(\vec{x}_1) \wedge \dots \wedge S_n(\vec{x}_n) \rightarrow (\exists \vec{y}) G_1(\vec{y}_1) \wedge \dots \wedge G_n(\vec{y}_n)$$

is equivalent to the GLAV expression $Q^S(\vec{x}) \subseteq Q^G(\vec{x})$ where

- $Q^S(\vec{x}) :- S_1(\vec{x}_1), \dots, S_n(\vec{x}_n)$
- $Q^G(\vec{x}) :- G_1(\vec{y}_1), \dots, G_n(\vec{y}_n)$

From

$S_1.\text{Movie}(i, t), S_1.\text{MovieDetails}(i, d, g, y) \subseteq \text{Movie}(t, d, \text{"comedy"}, y), y \geq 1970$

To

$S_1.\text{Movie}(i, t) \wedge S_1.\text{MovieDetails}(i, d, g, y) \rightarrow \text{Movie}(t, d, \text{"comedy"}, y) \wedge y \geq 1970$

Reformulation with TGD's can be done relatively straightforwardly with the
Inverse-Rules algorithm