

Relational Databases

A Review

Guillaume Raschia — Université de Nantes

Last update: April 8, 2013

[Source : J. Cheney, Univ. of Edinburgh]

[Source : L. Libkin, Univ. of Edinburgh]

Databases

What is a database?

- Any collection?
- A file system?
- A collection of relational tables?
- A bunch of XML documents?
- A multimedia digital library containing text, images, music, and sound recordings?

What differentiates a big pile of data from a database?

Databases

A database is...

- ...a collection of **structured data** (files, records, trees, key-values)
- ...that provides a **high-level interface** (query/update languages)
- ...and isolates **users** from low-level (and transient) implementation details (storage layout, query optimization, data structure traversals)

In other words, databases are **ADTs for managing large amounts of information**

Database theory

Is it all just about making SQL Server 0.3% faster ? No.

- Database technology is very important commercially, hence an alphabet soup of acronyms, buzzwords, standards, and fads
- But the underlying theory of databases is elegant and technology-independent
- In fact, database systems (in their current form) would not exist without theory

When Dinosaurs Ruled the Earth

- 1960s: Databases provided “network” (graph) or “hierarchical” (tree) data model. CODASYL, COBOL. No underlying theory or declarative semantics

When Dinosaurs Ruled the Earth

- 1960s: Databases provided “network” (graph) or “hierarchical” (tree) data model. CODASYL, COBOL. No underlying theory or declarative semantics
- 1970s: Codd proposes relational model: everything is a table, separate interface from implementation, etc.

When Dinosaurs Ruled the Earth

- 1960s: Databases provided “network” (graph) or “hierarchical” (tree) data model. CODASYL, COBOL. No underlying theory or declarative semantics
- 1970s: Codd proposes relational model: everything is a table, separate interface from implementation, etc.
- 1980s: Commercial success of RDBMS technology (Oracle, IBM, etc.)

When Dinosaurs Ruled the Earth

- 1960s: Databases provided “network” (graph) or “hierarchical” (tree) data model. CODASYL, COBOL. No underlying theory or declarative semantics
- 1970s: Codd proposes relational model: everything is a table, separate interface from implementation, etc.
- 1980s: Commercial success of RDBMS technology (Oracle, IBM, etc.)
- 1990s: Object oriented DB research popular, but fails to make significant commercial impact

When Dinosaurs Ruled the Earth

- 1960s: Databases provided “network” (graph) or “hierarchical” (tree) data model. CODASYL, COBOL. No underlying theory or declarative semantics
- 1970s: Codd proposes relational model: everything is a table, separate interface from implementation, etc.
- 1980s: Commercial success of RDBMS technology (Oracle, IBM, etc.)
- 1990s: Object oriented DB research popular, but fails to make significant commercial impact
- 2000s: XML/semistructured DB research popular, but...?

When Dinosaurs Ruled the Earth

- 1960s: Databases provided “network” (graph) or “hierarchical” (tree) data model. CODASYL, COBOL. No underlying theory or declarative semantics
- 1970s: Codd proposes relational model: everything is a table, separate interface from implementation, etc.
- 1980s: Commercial success of RDBMS technology (Oracle, IBM, etc.)
- 1990s: Object oriented DB research popular, but fails to make significant commercial impact
- 2000s: XML/semistructured DB research popular, but...?
- 2010s: NoSQL DB systems: cloud for scaling-up, simple makes it efficient, anti-relational

When Dinosaurs Ruled the Earth

- 1960s: Databases provided “network” (graph) or “hierarchical” (tree) data model. CODASYL, COBOL. No underlying theory or declarative semantics
- 1970s: Codd proposes relational model: everything is a table, separate interface from implementation, etc.
- 1980s: Commercial success of RDBMS technology (Oracle, IBM, etc.)
- 1990s: Object oriented DB research popular, but fails to make significant commercial impact
- 2000s: XML/semistructured DB research popular, but...?
- 2010s: NoSQL DB systems: cloud for scaling-up, simple makes it efficient, anti-relational
- The future: ACID/SQL revolution with NewSQL?

Preaching to the choir

- In retrospect, widely agreed that simple, clear, compelling data model is what made RDBMSs so successful
- Lack of same is what killed network, hierarchical, OO databases
- Some feeling that XML DB work has same problem

So: Theory is—at least—as important than practice

Source material

- For relational database theory:
 - Foundations of Databases, S. Abiteboul, R. Hull, and V. Vianu, Addison-Wesley 1995.
 - Elements of Finite Model Theory, L. Libkin, Springer-Verlag 2004.
- For database research foundations:
 - Readings in Database Systems (*The Red Book*), J. H. Hellerstein, and M. Stonebraker, 4th Edition, MIT Press 2005.
- For database systems:
 - Database Systems: The Complete Book, H. Garcia-Molina, J. D. Ullman, and J. Widom, 2nd Edition, Prentice Hall 2008.
 - Database Management Systems, R. Ramakrishnan, and J. Gehrke, 3rd Edition, McGraw-Hill Science, 2002.

Overview

It is all about...

The relational data model

- Data is stored in n -ary relations (i.e., tables)
- Data is described using “schema” (relations are typed)
- Data is manipulated using set-at-a-time relational operators (SQL)

Behind the scenes

Key theoretical issues

- Query language semantics and logical interpretation
- Descriptive power of query language features: conjunction, disjunction, negation, quantification
- Decision problems: query equivalence, constraint entailment
- Related area: finite model theory

The relational model

Some terminology

- **Domain**: countably infinite set D of individual data values (by convention, $D = \{a, b, c, d, \dots\}$)
- Rows, records or **tuples** : a sequence of data elements $\langle \vec{d} \rangle$.
- **Relations** or tables: finite sets of records (of the same width)
- Database (**instance**): collection of named relations
- **Schema**: a description of the structure of a record, table, or database

Schemas

- Relations “typed” by number of arguments: $R:n$ iff $R \subset D^n$
- The only base type is D , so number of arguments suffices. If $R:n$, define **arity** $\text{ary}(R) = n$
- Database instances I are “typed” by schemas \mathcal{R}

$$\mathcal{R} ::= \{R_1:m_1, \dots, R_n:m_n\}$$

- Example:

$$R:2, S:3$$

means “relation R has two fields, and S has three fields.”

Instances

An instance I (of a database schema \mathcal{R}) is a collection of finite relations matching \mathcal{R}

$$\left\{ R = \{\langle a, b \rangle, \langle b, c \rangle\}, S = \{\langle a, a, a \rangle, \langle a, b, b \rangle\} \right\}$$

is an instance of $\mathcal{R} = \{R:2, S:3\}$

- Order of rows doesn't matter, duplicates irrelevant
- $\text{adom}(R) \subset D$ is the **active domain** of R , i.e. the set of domain values that occur in R
- Tabular notation:

R		S		
a	b	a	a	a
b	c	a	b	b

Exercises 1/2

1. Definitions

Database, Table, Relation, Attribute, Tuple, Schema, Instance, Domain, Value, Active domain

2. True or False?

- i) A relation may have an infinite number of tuples.
- ii) Every relation follows a schema.
- iii) Attributes are typed.
- iv) $I \subset J$ implies that $\text{adom}(I) \subset \text{adom}(J)$.

Exercises 2/2

3. Is it a relation?

Which of the following relations conform to the schema $R:2$?

1. $R = \{\langle a, b \rangle, \langle c, b \rangle\}$
2. $R = \{\langle a, b \rangle, \langle c \rangle\}$
3. $R = \{\langle a, b \rangle, \langle b, a \rangle\}$
4. $R = \{\langle a, b \rangle, \langle a, b \rangle\}$ (abuse of set notation)
5. $R = \{\langle a, b \rangle, \langle c, \text{NULL} \rangle\}$
6. $R = \{\} = \emptyset$
7. $R = \{\langle \rangle\}$

4. Problem

How many different ways are there to draw a relation of schema $S:n$ within m tuples?

Core algebra

Expressions:

$$\begin{aligned}v &::= i \mid a \\ Q &::= \langle a \rangle \mid R \mid \sigma_{v=w}(Q) \mid \pi_{\vec{i}}(Q) \mid Q \times Q'\end{aligned}$$

- Index i , domain value a
- Singleton (1-tuple constant) $\langle a \rangle$
- Relation variables R
- **S**election σ
- **P**rojection π
- **C**ross product \times

This core algebra is so-called **SPC**

Selection

- $\sigma_{v=w}(Q)$ selects those rows $\langle \vec{d} \rangle$ satisfying $v = w$ from Q
- Here, v , w are either indices or domain constants
- Example:

$$\sigma_{1=2} \left(\begin{array}{|c|c|c|} \hline a & a & b \\ \hline a & b & c \\ \hline b & b & d \\ \hline b & c & c \\ \hline c & c & d \\ \hline \end{array} \right) = \begin{array}{|c|c|c|} \hline a & a & b \\ \hline b & b & d \\ \hline c & c & d \\ \hline \end{array}$$

Projection

- $\pi_{\vec{i}}(Q)$ projects the field indices listed in \vec{i} , in order (that is, drops all other fields)
- Example:

$$\pi_{1,3} \left(\begin{array}{|c|c|c|} \hline a & a & b \\ \hline a & b & b \\ \hline b & b & d \\ \hline b & c & c \\ \hline c & c & d \\ \hline \end{array} \right) = \begin{array}{|c|c|} \hline a & b \\ \hline b & d \\ \hline b & c \\ \hline c & d \\ \hline \end{array}$$

- Set-theoretic semantics: **duplicate elimination**

Cross product

- $Q \times Q'$ is the cross product of Q and Q'
- Example:

Q				Q'			$Q \times Q'$				
a	a	b	\times	a	55	$=$	a	a	b	a	55
a	b	c		c	42		a	a	b	c	42
b	b	d					a	b	c	a	55
							a	b	c	c	42
							b	b	d	a	55
							b	b	d	c	42

Normalization

- Two queries are **equivalent** ($Q_1 \equiv Q_2$) iff for every input, they produce the same output
- There are obviously lots of ways of writing “the same” query
- It would be handy to have a **canonical** representation of a given query

Theorem

Every SPC query has a normal form such that $Q \equiv Q'$ if and only if $\text{norm}(Q) = \text{norm}(Q')$

General idea: Rewrite to form

$$\pi_{\vec{J}}\left(\sigma_F(R_1 \times \cdots \times R_n)\right)$$

Normalizing selection and projection

- Neighboring projections can be composed

$$\pi_{\vec{j}}(\pi_{\vec{k}}(Q)) = \pi_{\vec{\ell}}(Q)$$

where $\ell_i = k_{j_i}$ for each i

- Problem: Neighboring selections commute

$$\sigma_{i=j}(\sigma_{i'=j'}(Q)) = \sigma_{i'=j'}(\sigma_{i=j}(Q))$$

- Solution: Allow **sets** (conjunction) of equations in selections

$$\sigma_F(\sigma_{F'}(Q)) = \sigma_{F \cup F'}(Q)$$

More rewriting rules

$$\sigma_F(\pi_{\vec{j}}(Q)) = \pi_{\vec{j}}(\sigma_{F'}(Q))$$

$$\text{where } F' = F[j_i/i]$$

$$\sigma_{1=j}(\langle a \rangle \times Q) = \langle a \rangle \times \sigma_{(j-1)=a}(Q)$$

$$((Q_1 \times \cdots \times Q_n) \times Q) = (Q_1 \times \cdots \times Q_n \times Q)$$

$$(Q \times (Q_1 \times \cdots \times Q_n)) = (Q \times Q_1 \times \cdots \times Q_n)$$

$$Q \times Q' = \pi_{\vec{\ell}\vec{m}}(Q' \times Q)$$

$$\text{where } \vec{m} = 1, \dots, \text{ary}(Q'),$$

$$\vec{\ell} = \text{ary}(Q') + 1, \dots, \text{ary}(Q) + \text{ary}(Q')$$

$$\sigma_F(Q) \times Q' = \sigma_F(Q \times Q')$$

$$\pi_{\vec{\ell}}(Q) \times Q' = \pi_{\vec{\ell}}(Q \times Q')$$

Numbers vs. names

- So far, we have used indices to refer to values in records
- This is convenient from a theoretical perspective because such expressions are concise and easy to specify
- But in real life (e.g. SQL), **field names** are useful
- Example. Guess what this does:

$$\pi_{1,2,5}(\sigma_{1=4}(R \times S))$$

How about this?

$$\pi_{\text{Name,Address,Phone}}(\sigma_{\text{Name}=P\text{Name}}(\text{AddressDir} \times \text{PhoneDir}))$$

Schemas with names

- Relations $R:\langle\vec{A}\rangle$ are typed by list of field names
- We say **sort** of R is $\langle\vec{A}\rangle$
- The only base type remains D , so list of names suffices
- Database instances are typed by schemas \mathcal{R} mapping relation names to sorts
- Example:

$$\mathcal{R} = \{R:\langle A, B \rangle, S:\langle A, B, C \rangle\}$$

means “relation R has two fields named A , B , and relation S has three fields named A , B , C .”

A running example

Cinema database

— Cinema = {Movie, Featuring, Location, Schedule} —

Movie: title, director, length, release_date

Featuring: title, actor, role

Location: theater, address, phone_number

Schedule: theater, title, showtime

Instances with names

- A named record is a finite map $\langle A_1:d_1, \dots, A_n:d_n \rangle$ from names to values
- By convention, field order doesn't matter
- A named relation $R:\langle \vec{A} \rangle$ is a set of named records
- An instance of \mathcal{R} is a collection of named relations matching named schema \mathcal{R}
- An instance of $\mathcal{R} = \{R:\langle A, B \rangle, S:\langle A, B, C \rangle\}$:

R

A	B
a	a
b	c
c	c

S

A	B	C
a	a	a
b	d	e
c	e	f

Core algebra with names

$$v ::= A \mid a$$

$$Q ::= \langle A:a \rangle \mid R \mid \sigma_{v=w}(Q) \mid \pi_{\vec{A}}(Q) \mid Q \bowtie Q' \mid \rho_{A_1 \dots A_n \rightarrow B_1 \dots B_n}(Q)$$

- Singleton constants $\langle A:a \rangle$, relation variables R
- **S**election σ
- **P**rojection π
- **J**oin \bowtie
- **R**enaming ρ

This algebra is called **SPJR**

Selection

- $\sigma_{v=w}(Q)$ selects those rows satisfying $v = w$ from Q
- Here, v , w are either constants or implicit field labels
- Example:

$$\sigma_{A=B} \left(\begin{array}{|c|c|c|} \hline A & B & C \\ \hline a & a & b \\ a & b & c \\ b & b & d \\ b & c & c \\ c & c & d \\ \hline \end{array} \right) = \begin{array}{|c|c|c|} \hline A & B & C \\ \hline a & a & b \\ b & b & d \\ c & c & d \\ \hline \end{array}$$

Projection

- $\pi_{\vec{A}}(Q)$ projects the field names listed in \vec{A} (that is, drops all other fields)
- Example:

$$\pi_{A,C} \left(\begin{array}{|c|c|c|} \hline A & B & C \\ \hline a & a & b \\ a & b & b \\ b & b & d \\ b & c & c \\ c & c & d \\ \hline \end{array} \right) = \begin{array}{|c|c|} \hline A & C \\ \hline a & b \\ b & d \\ b & c \\ c & d \\ \hline \end{array}$$

Join

- $Q \bowtie Q'$ joins Q and Q' by merging all pairs of records whose **common fields are equal**. Duplicate fields are omitted
- Example:

A	B	C
a	a	b
a	b	c
b	b	d
b	c	c
c	c	d

 \bowtie

A	D
a	55
a	100
c	42

 $=$

A	B	C	D
a	a	b	55
a	a	b	100
a	b	c	55
a	b	c	100
c	c	d	42

Renaming

- The renaming operator $\rho_{A_1 \dots A_n \rightarrow B_1 \dots B_n}(R)$ applies the simultaneous substitution $[B_1/A_1, \dots, B_n/A_n]$ to the field names of R .
- Example:

$$\rho_{BC \rightarrow DB} \left(\begin{array}{|c|c|c|} \hline A & B & C \\ \hline a & a & b \\ a & b & c \\ b & b & d \\ b & c & c \\ c & c & d \\ \hline \end{array} \right) = \begin{array}{|c|c|c|} \hline A & D & B \\ \hline a & a & b \\ a & b & c \\ b & b & d \\ b & c & c \\ c & c & d \\ \hline \end{array}$$

Equivalence

- Unnamed and named relations are equivalent
- Named [records, relations, instances, schemas] are in **bijective correspondence** with unnamed [records, relations, instances, schemas]
- Howto: fix an ordering on field names; assume all field name sequences are in increasing order

$$\langle d_1, \dots, d_n \rangle : n \Leftrightarrow \langle A_1 : d_1, \dots, A_n : d_n \rangle : \langle \vec{A} \rangle$$

Equivalence

- A SPJR query $Q:\langle \vec{A} \rangle$ on \mathcal{R} defines a function from named instances I of \mathcal{R} to named relations $Q(I)$ over $\langle \vec{A} \rangle$
- We say that two query languages are **equivalent** if their expressions define the same functions

Theorem

SPC and SPJR queries are equivalent in expressive power

- This is good because SPC is somewhat “lower level”: easier to analyze, but not as convenient
- Hence, can “compile” SPJR queries down to SPC queries without loss of expressiveness

Equivalence: sketch of the proof

It suffices to translate SPC queries to SPJR queries and *vice-versa*

- Selection, projection, renaming cases are easy
- **Key case 1:** $(R \bowtie S)^* = \pi_{\vec{i}}(\sigma_{i_1=j_1, \dots, i_n=j_n}(R^* \times S^*))$ where i_1, j_1 , etc. are indices of equal field names in R and S
- **Key case 2:** $(R^* \times S^*) = (\rho_{A_1 \dots A_n \rightarrow B_1 \dots B_n}(R)) \bowtie S$ where $\langle \vec{B} \rangle$ is all distinct from S field names
- Careful proof/implementation requires **bookkeeping** (to maintain mapping between field names and positions)

Exercises 1/2

1. Definitions

Algebra, SPC, SPJR, Query normal form, Query equivalence

2. True or False?

- i) Renaming is only “syntactic sugar” in SPJR.
- ii) Every SPC query is expressible within SPJR and *vice-versa*.
- iii) $\pi_{\vec{j}}(\pi_{\vec{k}}(R)) = \pi_{\vec{k}}(\pi_{\vec{j}}(R))$.
- iv) $\{\} \text{ op } R = R, \text{ op } \in \{\times, \bowtie\}$.
- v) $\{\langle \rangle\} \text{ op } R = R, \text{ op } \in \{\times, \bowtie\}$.
- vi) $\sigma_{i=a \vee j=b}(R)$ is expressible within SPC.
- vii) $\sigma_{i \neq a}(R)$ is not expressible within SPC.

Exercises 2/2

3. SPC and SPJR

1. Give the all schedule of **Katorza** theater.
2. Find theaters that show some movies directed by **Chabrol**.
3. Give the addresses of theaters that play some movies featuring **Jaoui** and where the director is also an actor.

4. Problem

Let \mathcal{R} be a database schema and Q a SPC query over \mathcal{R} ;

1. Prove that $Q(I)$ is finite for each instance I of \mathcal{R} .
2. Given instance I of \mathcal{R} and output arity n for SPC query $Q(I):n$, show an upper bound for the number of tuples that can occur in $Q(I)$. Show that this bound can be achieved.

Conjunctive calculus

- Calculus is a query formalism based on **set comprehension** notation
- Fragment of **first-order logic** (FO)

$$v ::= a \mid x$$

$$\Phi ::= R(\vec{v}) \mid v = v' \mid \exists x. \Phi \mid \Phi \wedge \Psi$$

$$Q ::= \{(\vec{v}) \mid \Phi\}$$

- Because only conjunctions are allowed in Φ , these are called **conjunctive queries**

Evaluation of formula

Semantics

Returns all tuples \vec{v} such that Φ is true

- **Free variables** $\text{free}(\Phi)$ are those that occur in \vec{v}
- $\text{adom}(\Phi)$ is the **active domain** of Φ , i.e. set of constants in Φ
- Valuation ν maps variables to constants from D : $\nu(\vec{x}) = \vec{a}$
- I satisfies Φ under ν denoted $I \models \Phi[\nu]$

$$Q(I) = \{\nu(\vec{v}) \mid I \models \Phi[\nu] \text{ and } \nu \text{ is a valuation over } \vec{v}\}$$

More examples

Find addresses of theaters that play movies directed by Chabrol

$$Q = \{(x, y) \mid \exists z. \text{Movie}(z, \textit{Chabrol}, -, -) \wedge \\ \textit{Location}(x, y, -) \wedge \textit{Schedule}(x, z, -)\}$$

Expressiveness

- Selection:

$$\sigma_{v=w}(R) = \{(\vec{x}) \mid R(\vec{x}) \wedge \vec{x}_{[v]} = \vec{x}_{[w]}\}$$

where $\vec{x}_{[i]} = x_i, \vec{x}_{[a]} = a$

- Projection:

$$\pi_{\vec{\ell}}(R) = \{(x_{\ell_1}, \dots, x_{\ell_k}) \mid R(\vec{x})\}$$

- Cross product:

$$R \times S = \{(\vec{x}, \vec{y}) \mid R(\vec{x}) \wedge S(\vec{y})\}$$

Tableau queries

- (\mathbf{T}, \vec{u}) : it consists of a set of tables, so-called **Tableau T** with constant and variable entries, and an answer row \vec{u}




Tableau query: example

Find addresses of theaters that play movies directed by Chabrol

Movie			
x_{ti}	Chabrol	x_{le}	x_{rd}


Location		
x_{th}	x_{ad}	x_{pn}

Schedule		
x_{th}	x_{ti}	x_{st}

$\Rightarrow \langle x_{th}, x_{ad} \rangle$


Expressiveness

- Selection $\sigma_{i=j}(R)$:



where $v_i = v_j = y$, $v_k = x_k$ otherwise


- Selection $\sigma_{i=a}(R)$:




where $v_i = a$, $v_k = x_k$ otherwise

Expressiveness (cont'd)

- Projection $\pi_{\vec{\ell}}(R)$:



- Cross product $R \times S$:



Conjunctive queries

So far, we've seen several **equivalent** query languages:

- SPC
- SPJR
- Conjunctive Calculus
- Tableau

These are all **conjunctive query** languages

Why is it so important?

Property 1

Conjunctive queries are **closed under composition**

For all CQs $Q' \vdash Q$ and $\mathcal{R} \vdash Q'$, then $Q \circ Q'$ is CQ as well

Property 2

Conjunctive queries are **satisfiable**

There exists an instance I of \mathcal{R} for which $Q(I)$ is non-empty

Property 3

Conjunctive queries are **monotonic**

For all instances I, J over \mathcal{R} such that $I \subseteq J$, then $Q(I) \subseteq Q(J)$

About satisfiability

- Conjunctive calculus with **equality** raises the following problem:

$$\{(\vec{x}) \mid R(\vec{x}) \wedge x_i = a \wedge x_i = b\} = Q^\emptyset$$

- **Unsatisfiable queries** could easily be checked by transitive closure of equality terms
- SPC and SPJR include unsatisfiable queries as well:
 - $\sigma_{\{i=a, i=b\}}(R) = Q^\emptyset$
 - $\sigma_{\{A=a, A=b\}}(R) = Q^\emptyset$

CQs are **satisfiable queries** only

Why is it so important? (cont'd)

Property 4

Conjunctive **query containment is decidable** and complexity is **NP-complete**

$\mathcal{R} \vdash Q \subseteq \mathcal{R} \vdash Q'$ iff for each instance I of \mathcal{R} , $Q(I) \subseteq Q'(I)$

- Small Q on large DB makes it acceptable
- Equivalence by checking mutual containment
 $Q \subseteq Q' \wedge Q \supseteq Q'$
- Very useful for query evaluation and optimization

Exercises 1/2

1. Definitions

Conjunctive query, Calculus, Tableau query, Satisfiability, Monotonicity

2. True or False?

- i) Conjunctive calculus is a fragment of FO.
- ii) SPC expressions are satisfiable queries only.
- iii) SPC and SPJR algebras are equivalent to CQs.
- iv) Conjunctive calculus w/o equality is no more CQ.
- v) NP-completeness of CQ query containment is a mess.
- vi) $\{R(1, 2)\}$ is a valid calculus formula.

Exercises 2/2

3. Conjunctive Calculus and Tableau Queries

1. Write calculus formulas for queries of the previous section.
2. Draw an unsatisfiable Tableau query.

4. Problem

1. Give evidence of satisfiability of CQs.
2. Give evidence of monotonicity of CQs.

What can't conjunctive queries handle?

Lots!

- Union \cup / disjunction \vee
- Difference $-$ / negation \neg and universal quantification \forall
- Recursive queries (e.g., transitive closure, connectivity)
- Primitive operations on D (arithmetic, string operations)
- Aggregation: counting, sum, average
- Complex data (nested records, variants, arbitrary trees)
- Arbitrary computations

Set operations

Relations are sets of tuples

So what about set operations?

- Intersection \cap
- Union \cup
- Complement or set difference $-$

Actually...

- Intersection is **already expressible** in SPC
- Union, complement are not
- How would one show this?
Idea: Follows from normal form theorem
- In fact, the value $\{\langle a \rangle, \langle b \rangle\}$ is not even expressible in SPC

Core algebra with union

$$v ::= i \mid a$$

$$Q ::= \langle a \rangle \mid R \mid \sigma_{v=w}(Q) \mid \pi_{\vec{i}}(Q) \mid Q \times Q' \mid Q \cup Q'$$

- **S**election σ
- **P**rojection π
- **C**ross product \times
- **U**nion \cup

This calculus is called **SPCU**

Named core algebra with union

$$\begin{aligned} v &::= A \mid a \\ Q &::= \langle A:a \rangle \mid R \mid \sigma_{v=w}(Q) \mid \pi_{\vec{A}}(Q) \\ &\quad \mid Q \bowtie Q' \mid \rho_{A_1 \dots A_n \rightarrow B_1 \dots B_n}(Q) \mid \mathbf{Q} \cup \mathbf{Q}' \end{aligned}$$

- **S**election σ
- **P**rojection π
- **J**oin \bowtie
- **R**enaming ρ
- **U**nion \cup

This calculus is called **SPJRU**

Difference

- Difference $Q - Q'$ can easily be added to [SPCU, SPJRU]
- [SPCU, SPJRU] + difference are called the (named)

Relational Algebra

- Union/disjunction and difference/negation pose little or no difficulty for algebra

Theorem

Query equivalence is **no more decidable** in RA

So far so good!

Add syntax for disjunction and negation to **relational calculus**:

$$\begin{aligned}\Phi &::= \exists x. \Phi \mid R(\vec{v}) \mid v = v' \mid \Phi \vee \Psi \mid \Phi \wedge \Psi \mid \neg \Phi \\ Q &::= \{(\vec{v}) \mid \Phi\}\end{aligned}$$

- Really, Φ can be an arbitrary first-order formula since $\forall x. \Phi \equiv \neg \exists x. \neg \Phi$
- Very really, \wedge, \exists, \neg suffice since $\Phi \vee \Psi \equiv \neg(\neg \Phi \wedge \neg \Psi)$
- **Semantics**: same as for ordinary first-order (FO) formulas?

Disjunction: problems

Conjunctive calculus + disjunction = **positive calculus**

Positive calculus is equivalent to [SPJRU, SPCU]

$$\{(x, y, z) \mid R(x, y) \vee R(y, z)\}$$

- For nonempty R , **answer is infinite** due to free x or z
- Unsafe query may produce infinite output from finite input
- Because databases are finite, we restrict to **safe queries**

Negation: problems

- Give it a try:

$$\{(x) \mid \neg R(x)\}$$

- Obviously unsafe. . . and this one:

$$\{(x) \mid \forall y. R(x, y)\} = \{(x) \mid \neg \exists y. \neg R(x, y)\}$$

- Whoa! This query is **domain-dependent**: answer depends on domain of quantification
 - e.g. if $R(a, b)$, $R(a, a)$ hold, and y ranges over $\{a, b\}$, then $x = a$ works. . .
 - but if y ranges over $\{a, b, c\}$ then $x = a$ doesn't work!

Safe domain-independent queries

Wanted: Guarantees that FO makes sense as a query language

- Unfortunately, **domain independence is undecidable**
- Two solutions:
 1. Find decidable sufficient conditions for domain dependence; prove that no loss of expressiveness ensues
 2. Observe that infinite models have finite descriptions in terms of equations; generalize database theory to finite sets of constraints rather than finite sets of facts
- Approach #2 is more general/satisfying; usually, focus on #1

Safety analysis

Range-restricted variable $x \in \text{free}(\Phi)$ s.t. $\nu(x) \in \text{adom}(\Phi, I)$

Basic idea: **every head variable must be constrained in the body**

$$\begin{aligned}rr(R(\vec{v})) &= \text{free}(R(\vec{v})) \\rr(x = a) = rr(a = x) &= \{x\} \\rr(\Phi \wedge x = y) &= \begin{cases} rr(\Phi) & (x, y \notin rr(\Phi)) \\ rr(\Phi) \cup \{x, y\} & (x \in rr(\Phi) \text{ or } y \in rr(\Phi)) \end{cases} \\rr(\Phi \wedge \Psi) &= rr(\Phi) \cup rr(\Psi) \\rr(\Phi \vee \Psi) &= rr(\Phi) \cap rr(\Psi) \\rr(\neg\Phi) &= \emptyset \\rr(\exists x. \Phi) &= \begin{cases} rr(\Phi) - \{x\} & (x \in rr(\Phi)) \\ \perp & (x \notin rr(\Phi)) \end{cases}\end{aligned}$$

$\{\vec{v} \mid \Phi\}$ is **safe-range** iff for some $\Psi \equiv \Phi$, then $\text{free}(\vec{v}) \subseteq rr(\Psi)$

Exercises 1/2

1. Definitions

Relational algebra, Positive calculus, Safe query,
Domain-dependent query, Safe-range query

2. True or False?

- i) Union and difference make CQ become RA.
- ii) Safe-range RC is equivalent to RA.
- iii) Query containment can be evaluated algorithmically.
- iv) It is possible to express unsafe query in RA.
- v) Domain-dependancy is decidable.
- vi) $\{(x, y, z) \mid R(x, y) \vee R(y, z)\}$ is finite when x, y, z in $\text{adom}(R)$.

Exercises 2/2

3. RA and RC

1. Which theaters do not show any movies directed by **Chabrol**?
2. Which theaters show only movies directed by **Chabrol**?
3. Which theaters show all movies directed by **Chabrol**?

4. Problem

For each of the following queries, guess whether it is domain-independent and/or safe range.

$$\{(x, y) \mid \exists z.(R(x, z) \wedge \exists w.S(w, x, y)) \wedge x = y\} \quad (1)$$

$$\{(\rangle \mid \exists x \forall y.(R(y) \rightarrow S(x, y))\} \quad (2)$$

$$\{(x, y) \mid (x = a \vee \exists z.R(y, z)) \wedge S(y)\} \quad (3)$$

If it is not domain-independent, exhibit a counter-example; and if it is safe range, translate it into an RA expression.

Declarative vs. Procedural

- Query languages are **declarative**: **what** in the output?

$$\{x_{th} \mid \exists x_{ti}. \text{Movie}(x_{ti}, \text{Chabrol}, -, -) \wedge \text{Schedule}(x_{th}, x_{ti}, -)\}$$

- Database system operates internally with different, **procedural** languages, which specify how to get the result

```
1 for each tuple T1=(ti1,di,le,rd) in table Movie do
2   for each tuple T2=(th,ti2,ts) in table Schedule do
3     if ti1=ti2 and di='Chabrol' then output th
4   end
5 end
```

Declarative vs. Procedural (cont'd)

Theoretical languages

- Declarative: relational calculus
- Procedural: relational algebra

Practical languages

Mix of both but mostly one use declarative features

- QBE (Tableau queries)
- SQL

SQL

- Structured Query Language
- Developed originally at IBM in the late 70s
- First standard: SQL-86 (with minor revision in SQL-89)
- Second standard: SQL-92
- Latest standard: SQL-99, or **SQL3**, well over 1,000 pages
- SQL2003 and SQL2008 add extra features
- De-facto standard of the relational database world—replaced all other languages

Example of SQL queries

Reminder of the database schema:

Movie: title, director, length, release_date

Featuring: title, actor, role

Location: theater, address, phone_number

Schedule: theater, title, showtime

Find titles of current movies

```
1 SELECT Title
2 FROM Movie ;
```

- SELECT lists attributes that go into **the output** of a query
- FROM lists **input relations**
- Algebraic formula: $\pi_{\text{title}}(\text{Movie})$
- **Warning:** SQL uses **bags** rather than sets

Relational operations on bags

RDBMS implements relations as bags (or multisets)

- Union adds occurrences: $|R \cup S| = |R| + |S|$
- Duplicates are preserved in projection: $|\pi_{\vec{A}}(R)| = |R|$
- Consistent aggregate computation: $\text{AVG}(R[\vec{A}])$
- t in $R \cap S$: $\min(|R(t)|, |S(t)|)$
- t in $R - S$: $\max(0, |R(t)| - |S(t)|)$
- No surprise for σ , \times , \bowtie
- Algebraic laws revisited: $(R \cup S) - T = (R - T) \cup (S - T)$?

Relations as multisets

The results for Conjunctive Query containment over relations-as-sets no longer hold!

Property 1

CQ containment is **no longer in NP**

Property 2

Query containment of unions of CQ is **undecidable**

Lesson:

Small changes in the data model can have a big impact

Example of SQL queries (cont'd)

—— Find theaters showing movies directed by Chabrol ——

```
1 SELECT S.Theater
2 FROM Schedule S, Movie M
3 WHERE M.Title = S.Title
4        AND M.Director='Chabrol' ;
```

New features:

- SELECT now specifies which relation the attributes came from—because we use more than one
- FROM lists two relations with aliases
- WHERE specifies the **condition** for selecting a tuple

Translation to SPC

Conjunctive Query

SELECT DISTINCT-FROM-WHERE SQL queries with conjunct of equality conditions are equivalent to SPC

- Algebraic expression:

$$\underbrace{\pi_{S.Theater}}_{\text{SELECT statement}} \left(\underbrace{\sigma_{M.Title=S.Title \wedge M.Director="Chabrol"}}_{\text{WHERE statement}} \left(\underbrace{\text{Schedule } S \times \text{Movie } M}_{\text{FROM statement}} \right) \right)$$

SQL and RC

Tuple Relational Calculus

Core declarative part of SQL is close to TRC, a variant of Relational Calculus

$$\{t.theater \mid \exists u.Movie(u) \wedge Schedule(t) \wedge \\ t.title = u.title \wedge u.director = 'Chabrol'\}$$

- Same operators than (Domain)RC
- Variables are **tuples**
- Attribute values can be reached b.t.w. of dot notation

TRC is obviously equivalent to (D)RC

Joining relations

- WHERE allows to join together several relations
R.Title = S.Title
- A JOIN statement also exists in SQL

```
1  SELECT S.Theater
2  FROM Movie M NATURAL JOIN Schedule S
3  WHERE M.Director='Chabrol' ;
```

Join variants

Natural join is the \bowtie operator from SPJR

- Preferred SQL statement:

```
1  SELECT S.Theater FROM Movie M
2     INNER JOIN Schedule S USING (Title)
3     WHERE M.Director='Chabrol' ;
```

- Family of join operators in SQL:
 - **Equi-join**: $R \text{ INNER JOIN } S \text{ ON } R.A = S.B$
 - **θ -join**: $R \text{ INNER JOIN } S \text{ ON } R.A > S.B$
 - **Outer join**: $R \text{ LEFT OUTER JOIN } S \text{ ON } R.A \neq S.B$
 - **Cross product**: $R \text{ CROSS JOIN } S$

Positive Relational Algebra within SQL

Reminder

SPJR+Union operator makes Positive RA

Find actors who played in movies directed by Chabrol *OR* Polanski

```
1 SELECT F.Actor FROM Featuring F JOIN Movie M USING (Title)
2   WHERE M.Director='Chabrol' OR M.Director='Polanski' ;
```

SQL has also a dedicated UNION operator

```
1   (SELECT Actor FROM Featuring JOIN Movie USING (Title)
2    WHERE Director='Chabrol')
3 UNION
4   (SELECT Actor FROM Featuring JOIN Movie USING (Title)
5    WHERE Director='Polanski') ;
```

More Thoughts on Union

- UNION ALL to **preserve duplicates**
- **Renaming** to the rescue

_____ List all directors or actors _____

```
1      (SELECT Director AS Person FROM Movie)
2  UNION
3      (SELECT Actor AS Person FROM Featuring) ;
```

Intersection and Difference

\cap is expressible within SPJRU, $-$ is not

$$R \cap S = \rho_{\text{sort}(R) \rightarrow \vec{X}}(R) \bowtie \rho_{\text{sort}(S) \rightarrow \vec{X}}(S)$$

SQL syntax for

- $R \cap S$: R INTERSECT S or R INTERSECT ALL S
- $R - S$: R EXCEPT S or R EXCEPT ALL S

Find all actors who are...

_____ not directors _____

```
1 (SELECT Actor AS Person
2   FROM Featuring)
3 EXCEPT
4 (SELECT Director AS Person
5   FROM Movie) ;
```

_____ also directors _____

```
(SELECT Actor AS Person
  FROM Featuring)
INTERSECT
(SELECT Director AS Person
  FROM Movie) ;
```

Beyond simple queries

- So far we mostly translated $RA = \{\sigma, \pi, \bowtie, \rho, \cup, -\}$ into SQL
- Other SQL statements allow to express complex queries in a slightly different way than pure RA
 - Nested queries
 - “For all” queries
- Also, extra SQL features go far beyond RA

Nested queries

- WHERE clause could contain **subquery**

— Find actors who did not play in a movie by Chabrol —

```
1 SELECT F.Actor FROM Featuring F
2 WHERE F.Actor NOT IN (SELECT F1.Actor FROM Featuring F1
3                        JOIN Movie M USING (Title)
4                        WHERE M.Director='Chabrol') ;
```

- Nested query could be **correlated** to the outermost query

— Find dirs. whose movies are playing at Le Katorza —

```
1 SELECT M.Director FROM Movie M
2 WHERE EXISTS (SELECT * FROM Schedule S JOIN M USING (Title)
3              WHERE S.Theater='Le Katorza') ;
```

- Actually, nested query could occur anywhere!

For all in SQL

Find directors whose movies are playing in all theaters

$$\{(x.\text{director}) \mid M(x) \wedge \forall y[L(y), \exists z(S(z) \wedge \\ z.\text{title} = x.\text{title} \wedge y.\text{theater} = z.\text{theater})]]\}$$

- M stands for Movie and S for Schedule and L for Location

$$\pi_{\text{director}}(M) - \pi_{\text{director}}\left(\left(\pi_{\text{theater}}(L) \times \pi_{\text{director}}(M)\right) - \pi_{\text{theater,director}}(M \bowtie S)\right)$$

- RA query is much less intuitive than RC query

For all in SQL (cont'd)

SQL's way of saying this

Find directors such that **there does not exist** a theater where their movies **do not** play

- Main idea: $\forall x.P(x) \Leftrightarrow \neg \exists x.\neg P(x)$

[illegible]

Relational division

$$\alpha \div \beta = \pi_{A-B}(\alpha) - \pi_{A-B}\left(\left(\pi_{A-B}(\alpha) \times \beta\right) - \alpha\right)$$

At least four ways to write division in SQL

1. Direct conversion of the relational algebra expression:
translate $-$, π , \times
2. Logical tautology (see previous slide)
3. Set containment: if $X \supseteq Y$, then $Y - X = \emptyset$ (or $\neg \exists x. x \in Y - X$)
4. Set cardinalities: check for $|Y| - |X| = 0$. Need for SQL group by/having and count, discard negation

Other features of SQL

- Datatypes, type-specific operations
- Table declaration, constraint enforcement (DDL part)
- Database modifications: insert, update, delete
- NULL values
- Views and temporary tables
- Aggregation

Simple aggregate queries

Count the number of Movies

```
1      SELECT COUNT(*)  
2      FROM Movie ;
```

Add up all movie lengths

```
1      SELECT SUM(Length)  
2      FROM Movie ;
```

Find the number of directors

```
1      SELECT COUNT(DISTINCT Director)  
2      FROM Movie ;
```

Aggregation and grouping

For each theater playing at least one long (over 2 hours) movie, find the average length of all movies played there

```
1 SELECT S.Theater, AVG(M.Length) AS Average_Length
2 FROM Schedule S
3 JOIN Movie M USING (Title)
4     GROUP BY S.Theater
5     HAVING MAX(M.Length) > 120
```

Exercises 1/2

1. Definitions

SQL, Bag-RA, Nested query, θ -join, Outer join

2. True or False?

- i) Every RA query is expressible within SQL.
- ii) $\text{SELECT } R.A \text{ FROM } R \equiv \pi_A(R)$.
- iii) $\{(1), (1), (2)\} \text{ EXCEPT } \{(2), (3)\} = \{(1), (1)\}$.
- iv) $(R \cup S) - T = (R - T) \cup (S - T)$ with SQL bag semantics.
- v) Boolean expression with a NULL value returns True or False.

Exercises 2/2

3. SQL queries

1. Find actors who did not play in a movie by **Chabrol**.
2. Find theaters that play movies not played anywhere else.
3. Write division in the *translated RA* flavor.

4. Problem

"For each theater playing at least one long (over 2 hours) movie, find the average length of all movies played there."

The above aggregate query has a straightforward SQL translation with GROUP BY/HAVING statement. Show another way to write SQL query without GROUP BY/HAVING. *Hint*: use subqueries in the SELECT clause and in the WHERE clause.

Database constraints

- In our running examples we assumed that the title attribute identifies a movie
- But this may not be the case:

title	director	length	release_date
Dracula	Browning	84mn	1931
Dracula	Fischer	82mn	1958
Dracula	Badham	109mn	1979
Dracula	Coppola	127mn	1992

- Database constraints: provide additional semantic information about the data
- Most common ones: **functional** and **inclusion** dependencies, and their special cases: **keys** and **foreign keys**

Functional dependency

- If we want the title to identify a movie uniquely (i.e., no multiple Dracula records), we express it as a **functional dependency**

title \longrightarrow director, length, release_date

- More generally:

$$X \longrightarrow Y ::= \left(\forall t, u \in R, t[X] = u[X] \implies t[Y] = u[Y] \right)$$

Running Example

R

A	B	C	D	E
a	a	b	a	d
a	b	b	a	d
a	c	c	d	d
b	a	a	a	d
b	b	a	a	d
b	d	d	c	a

Functional dependencies that hold in R are:

$$\mathcal{F} = \{AB \longrightarrow CD, C \longrightarrow ADE, B \longrightarrow DE, D \longrightarrow E\}$$

Keys

Let K be a subset of attributes of $R:\langle \vec{U} \rangle$. Then K is a **key** if R satisfies functional dependency $K \longrightarrow U$ and K is minimal

Follow-on of the example

Keys are $\{AB, BC\}$

Among **candidate keys**, one can be arbitrarily promoted into **primary key**

Inclusion constraints

Referential integrity

Attributes of one relation refer to values in another one

- These particular constraints are called **inclusion dependencies** (ID)
- Formally, we have an inclusion dependency $R[X] \subseteq S[Y]$ when every value of the set of attributes X in R also occurs as a value of the set of attributes Y in S :

$$\pi_{\vec{X}}(R) \subseteq \pi_{\vec{Y}}(S)$$

Foreign keys

- Most often IDs occur as part of a **foreign key**
- Foreign key is a conjunction of a key and an ID:

$$R[X] \subseteq S[Y] \text{ and } Y \longrightarrow U, \text{ with } S:\langle \vec{U} \rangle$$

Example

We expect Theater and Title from Schedule to be found resp. in Location and Movie:

- $\text{Schedule}[\text{Theater}] \subseteq \text{Location}[\text{Theater}]$
- $\text{Schedule}[\text{Title}] \subseteq \text{Movie}[\text{Title}]$

If Title is a key for Movie, then it is **foreign key** in Schedule
Same arises for Theater

Reasoning about FDs

Closure

Denote by \mathcal{F} the set of functional dependencies on R ; **Closure** of \mathcal{F} is $\mathcal{F}^+ = \{f \mid \mathcal{F} \models f\}$

Attribute closure of X on \mathcal{F} : $X_{\mathcal{F}}^+ = \{A \mid X \longrightarrow A \in \mathcal{F}^+\}$

Armstrong's axioms

- **Reflexivity**: if $X \supseteq Y$, then $X \longrightarrow Y$
- **Augmentation**: if $X \longrightarrow Y$, then $XZ \longrightarrow YZ$ for any Z
- **Transitivity**: if $X \longrightarrow Y$ and $Y \longrightarrow Z$, then $X \longrightarrow Z$
- These are **sound** and **complete** inference rules for FDs!

Reasoning about FDs (cont'd)

Commonly derived rules:

- **Union:** if $X \longrightarrow Y$ and $X \longrightarrow Z$, then $X \longrightarrow YZ$
- **Decomposition:** if $X \longrightarrow YZ$, then $X \longrightarrow Y$ and $X \longrightarrow Z$
- **Pseudo-transitivity:** if $X \longrightarrow Y$ and $YZ \longrightarrow T$, then $XZ \longrightarrow T$

Back to the example

$$\mathcal{F}^+ = \mathcal{F} \cup \{A \longrightarrow A, AB \longrightarrow A, BC \longrightarrow D, \\ BC \longrightarrow BC, DA \longrightarrow E, AB \longrightarrow DE, ABCD \longrightarrow A, \\ ABCD \longrightarrow ABCD, \dots\}$$

Canonical cover

On the other side

Canonical cover \mathcal{F}_{\min} such that $\mathcal{F}_{\min}^+ = \mathcal{F}^+$ and FDs in \mathcal{F}_{\min} are all **irreducible**

- Non unique
- Preferred form for normalization

Example

$$\mathcal{F}_{\min} = \{AB \longrightarrow C, C \longrightarrow D, C \longrightarrow A, B \longrightarrow D, D \longrightarrow E\}$$

Idea: Is $AB \longrightarrow D$ redundant in \mathcal{F} ?

Check for $D \in AB_{\mathcal{G}}^+$, $\mathcal{G} = \mathcal{F} - \{AB \longrightarrow D\}$

Database design

Normalization

Avoid **redundancy** and modification **anomalies**

Example

Assume AC , BD and DE are three distinct entities in R

1. d value **occurs 5 times** in column E of R
2. **Updating** $E:a$ to $E:a'$ in the first row requires to update 4 more rows, making the database inconsistent otherwise
3. **Inserting** $\langle A:b, C:b \rangle$ in R requires to provide BDE values as well
4. **Deleting** last pair $\langle D:a, E:d \rangle$ in R implies to delete ABC values as well

Normal forms

1NF

Table has (a) a **key** and (b) **atomic** columns and (c) **no repeating groups** of columns

Examples

- R is 1NF
- Sets or tuples or tables are not allowed as attribute values
- $(\text{Author}_1, \text{Author}_2)$ is not allowed as a subset of columns

Normal forms (cont'd)

2NF

1NF + **full** FD from keys to non-prime attributes

Prime attributes are those that belong to any candidate key

Examples

- R is not 2NF since $B \longrightarrow D$ holds in R and AB is a key and D is a non-prime attribute
- $R_1 = \pi_{BDE}(R)$ and $R_2 = \pi_{ABC}(R)$ are both 2NF

Normal forms (cont'd)

3NF

2NF + **non-transitive** FD from keys to non-prime attributes

Examples

- $R_2: \langle ABC \rangle$ is 3NF with key AB
- $R_1: \langle BDE \rangle$ is not since $D \longrightarrow E$ holds
- $R_{11} = \pi_{BD}(R_1)$ and $R_{12} = \pi_{DE}(R_1)$ are both 2NF

Normal forms (cont'd)

BCNF

3NF + every non-trivial FD is on a **superkey**

Examples

- $R_{11}:\langle BD \rangle$ and $R_{12}:\langle DE \rangle$ are both BCNF
- $R_2:\langle ABC \rangle$ is not since $C \longrightarrow A$ holds
- $R_{21} = \pi_{AC}(R_2)$ and $R_{22} = \pi_{BC}(R_2)$ are both BCNF

Decomposing a relation

Decomposition of R into $S:\langle\vec{X}\rangle$ and $T:\langle\vec{Y}\rangle$

- Lossless-join decomposition of R : $\pi_{\vec{X}}(R) \bowtie \pi_{\vec{Y}}(R) = R$
- Lossless w.r.t. \mathcal{F} :

$$X \cap Y \longrightarrow X, \quad \text{or}$$

$$X \cap Y \longrightarrow Y$$

- Dependency preservation: $(\mathcal{F}_X \cup \mathcal{F}_Y)^+ = \mathcal{F}^+$

Theorem

*Lossless-join dependency preserving decomposition of R into a collection of 3NF relations is **always possible***

What about decomposition into BCNF?

Multivalued and join dependencies

More stringent normal forms

- MVD: 4NF...
- FD+JD: ETNF (H. Darwen et al., 2012)...
- JD: 5NF...
- JD: 6NF...
- Domain and Key constraints: DKNF...

Exercises 1/2

1. Definitions

Functional dependency, Inclusion dependency, Key, Foreign key, Closure, Armstrong's axioms, Canonical cover, Modification anomaly, 1NF, 2NF, 3NF, BCNF

2. True or False?

- i) A relation may have several keys.
- ii) Closure is not unique.
- iii) $\mathcal{F}_{\min} \subseteq \mathcal{F} \subseteq \mathcal{F}^+$.
- iv) $\emptyset \longrightarrow X$ always holds.
- v) $2^n \leq |\mathcal{F}^+| \leq 2^{2^n}$, where $n = \text{ary}(R)$.

Exercises 2/2

3. Misc.

Consider a database with a single relation $R:\langle abcde \rangle$; the set of functional dependencies that hold in R is

$$\mathcal{F} = \{ab \rightarrow cd, a \rightarrow d, e \rightarrow b, cd \rightarrow ce, ac \rightarrow bde, c \rightarrow a\}.$$

1. Give (i) keys of R , (ii) NF of R , (iii) a canonical cover of \mathcal{F} .
2. Decompose R up to the BCNF.

4. Problem — Armstrong Relations

- Prove that there exists an instance of $R:\langle \vec{U} \rangle$ such that for each FD f over U , $R \models f$ iff $f \in \mathcal{F}^+$. It is called an **Armstrong relation**.
- Exhibit a set \mathcal{F}_2 of FDs over $\{A, B, C\}$ such that each Armstrong relation for \mathcal{F}_2 has at least 4 distinct values occurring in the A column.

To Sum-up

1. The relational model is a set-theoretic approach to database
2. Theoretical languages: RA (procedural) \equiv RC (declarative)
3. Two flavors of RA: named (SPJRUD) or unnamed (SPCUD)
4. Nice fragments of RA: SPJR \equiv Conjunctive Queries
5. Practical languages: QBE (Tableau-like queries) and SQL (TRC-like queries at the heart)
6. Many sophisticated statements and extra-features in SQL
7. Functional and incl. dependencies between columns and tables
8. DB design prevents from redundancy and modification anomalies
9. Lossless-join dependency preserving decomposition up to 3NF

Database Management Systems

Main topics are

- **Physical model:** storage, access methods
- **Query optimization:** execution plan, cost evaluation, join algorithms, external sort
- **Transactions and concurrency control:** ACID properties, serializability, 2PL
- **Failure recovery:** logs
- **Tuning:** de-normalization, query tricks, administration

Great opportunities for self-learning!